

# DiCE: Online Testing of Deployed Federated and Heterogeneous Distributed Systems

EPFL Technical Report EPFL-REPORT-164475

*Marco Canini, Vojin Jovanović, Daniele Venzano, Gautam Kumar, Dejan Novaković, and Dejan Kostić*  
School of Computer and Communication Sciences, EPFL, Switzerland  
*email: firstname.lastname@epfl.ch*

## Abstract

It is notoriously difficult to make distributed systems reliable. This becomes even harder in the case of the widely-deployed systems that are heterogeneous (multiple implementations) and federated (multiple administrative entities). The set of routers in charge of the Internet’s inter-domain routing is a prime example of such a system. In this paper, we argue that a key step in making these systems reliable is the need to automatically explore the system behavior to check for potential faults. We present the design and implementation of DiCE, a system for online testing of heterogeneous and federated distributed systems. DiCE runs concurrently with the production system by leveraging distributed checkpoints and isolated communication channels. DiCE orchestrates the exploration of relevant system states by controlling the inputs that drive system actions. While respecting privacy among different administrative entities, DiCE detects faults by checking for violations of properties that capture the desired system behavior. We demonstrate the ease of integrating DiCE with a BGP router and a DNS server, the building blocks of two vital services in the Internet. Our evaluation in the testbed shows that DiCE quickly and successfully detects three important classes of faults, resulting from configuration mistakes, policy conflicts and programming errors.

## 1 Introduction

Many successful distributed systems are inherently heterogeneous and federated — heterogeneity arises from the creation of multiple, inter-operable implementations; federated refers to the existence of multiple service providers operating under different administrative domains. The Internet’s inter-domain routing system governed by BGP is a prime example of a heterogeneous and federated system. Other such systems include DNS, electronic mail, peer-to-peer content distribution [17],

content and resource peering [10], computing grids, and Web services. However, the resulting competing environment of mutually mistrusting providers fosters a tension between a provider’s own goals versus the common overarching desire of keeping the federated system functioning properly.

In such an environment, making distributed systems reliable does not stop with the already difficult task of producing a robust design and implementation. Achieving high reliability also bears the difficulties in deploying and operating these systems whose aggregate behavior is the result of interleaved actions of multiple system nodes running in a heterogeneous and failure-prone environment. In fact, several factors such as subtle differences in the details of inter-operable implementations, or system-wide conflicts due to locally admissible (mis)configurations can cause harmful node interactions that lead to faults, i.e., *deviations of system components from their expected behavior*. These faults which span the state and configuration across multiple nodes are perhaps less frequent than single-machine bugs, e.g., memory-related issues. However, when these faults manifest themselves they have far-reaching and substantial negative impact, and require considerable resources to be diagnosed and eliminated.

For example, a BGP router can rightfully decide to reset its peering session in response to a syntactically valid, but semantically ambiguous message. However, when many of such routers are coupled with another large number of routers that propagate the ambiguous message (because of a different message parser implementation), the overall effect is a large fraction of routers that are continuously resetting and restoring sessions as it happened in several episodes [6]. The resulting high update processing rate causes a performance and reliability problem. Others have argued that a malformed packet could take down a significant fraction of the Internet [3]. Even with a 100% protocol-compliant message, such an incident inadvertently occurred in August of 2010 [5].

Our overarching vision is to harness the continuous

increases in available computational power and bandwidth to improve the reliability of distributed systems. We argue that nodes in the system and their administrators should be proactively working towards finding which node actions could potentially lead to faults. This task cannot be done only locally by checking the single-node behavior, as the erroneous system state can span multiple nodes and remote node configurations are not available locally. Thus, detecting faults in the general case requires some collaboration among the nodes. The faults these actions lead to are evidence of possible future system failures which may be avoided by detecting these potential faults.

To detect faults, we propose to continuously and automatically explore the system behavior alongside the production system, but in complete isolation from it using a system snapshot captured from the current state. That is, we check system-wide consequences of a large number of actions nodes can undertake, and output actions that lead to failures. In practice, node actions are the result of subjecting the node’s code in its current state to messages, configuration changes, failures, random choices, etc., collectively called *inputs* in the following. Therefore, we want to subject nodes to a large number of possible inputs that systematically exercise their code paths.

We have to address several difficult challenges [14] in our work. The federated nature of many deployed systems means that a node cannot gain unrestricted access to other nodes’ state and configuration. Moreover, we have to carefully manage the information flowing between system participants to preserve their confidential nature. The heterogeneity of the system makes it difficult or impossible to have local access to the source or binary code of other participants. Systematically exploring node behavior even for a single node easily runs into the problem of exponential explosion in the number of code paths that need to be explored. Finally, the sheer size of the system can pose scalability problems.

Static analysis of configuration files [20] cannot be applied to this problem because it does not take into account the actual state and software of the system. Tools for predicting inconsistencies using live model checking (e.g., [42]) cannot be used because they require a node to (i) retrieve checkpoints (with private state and configuration) from other participants, and (ii) obtain access to the source code of other participants. Applying systematic source code exploration tools based on symbolic execution from initial state [13, 38] cannot explore code paths sufficiently deep due to exponential growth in the number of possible paths caused by having large inputs (configuration and messages received over a long time).

In this paper, we introduce DiCE, a system for on-line testing of heterogeneous, federated distributed systems. Accounting for the federated nature of the system,

we let each node autonomously explore its local actions. To exercise possible node actions, we use a technique called *concolic execution* [12, 18, 24] to produce the inputs that systematically explore all possible code paths at one node. We overcome the problem of exponential explosion of code paths by starting exploring the node behaviors from current system state, and by subjecting the node’s code to small-sized inputs that affect localized parts of state-changing code. We use a set of lightweight node checkpoints to allow the single node’s actions reach out to other nodes as a way to drive and explore system-wide state in isolation from the production environment. To preserve privacy between different administrative domains, we define a narrow information sharing interface that enables a node to query remote nodes for relevant state checks. We detect faults by checking and flagging violations of given properties that, tying together state checks over multiple system nodes, capture the desired system behavior.

We demonstrate DiCE’s ability to detect faults in two systems that provide fundamental services: BGP and DNS. The faults that DiCE detects are a result of long-standing, hard problems. For example, we are not patching particular BGP problems. Instead, we demonstrate the benefits of having a generic framework that enables system operators to specify the desired behavior and learn about possible faults and their impact.

DiCE is not a bug-finding tool per se — property violations may uncover several insidious misbehaviors due to harmful node interactions, the root cause of which includes configuration mistakes and programming errors.

DiCE is a crucial step in being able to guard against important classes of faults. Advance warnings can be used to notify the system operator(s) about a particular misconfiguration, or to trigger automatic or semi-automatic installation of a filter against the problem caused by the software reaction to an unanticipated message. A particular benefit of our approach is that the separate administrative entities can use DiCE by integrating only their source code with it, and without requiring access to the source code, executable, or configuration of other participants.

The contributions of this paper are as follows:

1. We describe the design and implementation of DiCE, a system for detecting possible faults in heterogeneous, federated environments.
2. We provide a technique for automatic and lightweight distributed snapshot creation that (i) respects trust boundaries, and (ii) allows system behavior to be explored in isolation. This technique effectively enables concolic execution to drive node actions and extend their reach across the network to explore relevant system state. We believe that this primitive can be successfully applied to other “what-if” exploratory scenarios.

3. We demonstrate how a small amount of input-producing code can be used to drive concolic execution across the relevant federated distributed system states. Doing so uncovers faults due to events that are difficult to explore, such as remote node failures. To the best of our knowledge, this is the first such approach.

4. We integrate DiCE with the BIRD [7] open-source router written in C. In our evaluation on the network testbed with Internet-like conditions, we demonstrate that DiCE quickly detects two important classes of faults that have affected the Internet: (i) Internet-wide BGP session resets, and (ii) policy conflicts among ISPs.

5. We integrate DiCE with the MaraDNS [4] open-source DNS server, and demonstrate its ability to detect cyclic zone dependencies – an insidious type of DNS misconfiguration that can render entire domain names unresolvable [36].

## 2 Design

To detect faulty states (those in which the system components deviate from their desired behavior), our goal is to continuously and systematically explore the behavior of a distributed system. In this section, we first offer an overview of how DiCE meets this goal. We then discuss each aspect of our design in detail, together with its rationale and principles.

**Overview** DiCE runs online, alongside a deployed system, off the critical execution path. Figure 1 gives a high-level illustration of how DiCE tests running distributed systems. First, one node in the system acts as an *explorer* (node marked with a double ellipse in step 1 of Figure 1). The explorer triggers the creation of a *shadow snapshot*, i.e., a consistent and distributed snapshot composed of nodes’ local checkpoints based on the current state of the system (step 2). Then, DiCE exercises a variety of local behaviors at the explorer (steps 3-5, stylized under the explorer node) that result in exploring system-wide relevant states. As detailed later, the code, current state, and inputs fed to the code determine how a node behaves. Thus, DiCE uses a combination of techniques to carefully construct the inputs that systematically explore node behavior. The execution of each node behavior occurs in isolation, over a clone of the shadow snapshot. The messaging is also confined to the cloned snapshot. Each cloned snapshot represents one instance of possible system behavior involving multiple nodes. DiCE detects faults by checking for violations of given safety and/or liveness properties in each cloned snapshot.

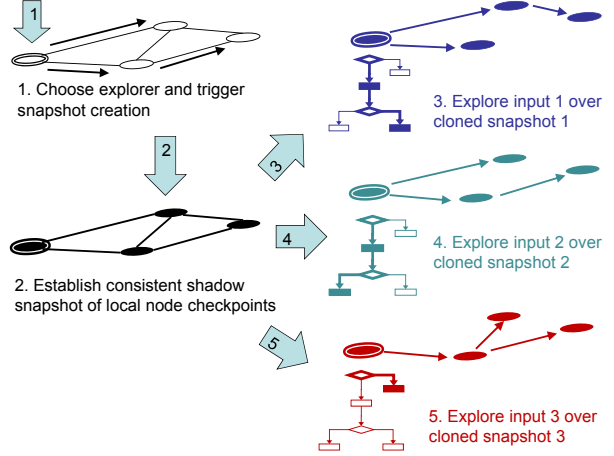


Figure 1: DiCE systematically explores and checks system behavior over isolated shadow snapshots.

### 2.1 Checkpointing state across nodes

Despite all the best efforts in thorough local testing and configuration checking, there is no substitute for having the ability to inspect distributed system state for potential faults. This is challenging because the federated nature of the systems we target makes it impossible to simply retrieve state checkpoints from other nodes. Moreover, it may be impossible to have the exact copy of the software running at other nodes, as the system is fundamentally heterogeneous. Finally, the entities controlling different nodes might not be willing to reveal their configurations.

**A snapshot that respects trust boundaries** Presented with these constraints, we decide to let nodes keep their state, code and configuration in a local checkpoint. A checkpoint never leaves the node that creates it. However, the checkpoint has the ability to clone itself and resume execution from its saved state and to communicate with checkpoints belonging to other nodes. This way, a node that wishes to explore its behavior can do so by creating and executing inside a shadow snapshot, i.e., a consistent, distributed set of individual node checkpoints of the explorer’s neighborhood.

The explorer establishes a snapshot by sending an annotated message to its immediate neighbors, which forward the message further on to their neighbors, etc. until a desired scope is reached.

**Isolating execution in a snapshot** To prevent the exploratory executions from changing system state, each node checkpoint is isolated from its environment. For example, all outgoing messages are intercepted and, instead of being transmitted over existing connections, are sent over *shadow connections* that the checkpointed node creates to the message destinations.

A node in a scalable distributed system typically interacts with its immediate neighbors. For good perfor-

mance, these neighbors are often picked to be in close proximity (in terms of network latency). Thus, exploration across the cloned snapshot should execute quickly in a limited scope.

## 2.2 Exploring system states

The key step in detecting potential faults is to explore a large number of possible system behaviors. In practice, the aggregate behavior is the result of interleaved actions of multiple system nodes. To explore system behavior under different scenarios, we could take a position atop the system from where we would control all individual node actions and their interleaving<sup>1</sup>. Unfortunately, this principle would create the need for a third party responsible for orchestrating state exploration in the targeted federated system. Also, when considering a large-scale system, several scalability issues would arise.

Because we want to let the nodes (and administrative domains) maintain control of how they participate in the system state exploration, we propose a different principle — focus on local actions of one node and let the exploration of a single node’s behavior reach out to other nodes as a way to explore system state. This kind of exploration can take place one node at a time, in parallel, or a combination thereof.

**What drives node behavior** In practice, the behavior of each node is determined by the path taken through its code. Keeping in mind that we resume execution from a local node checkpoint, we note that the code that will execute next is affected by (i) the current state and (ii) what we collectively term as inputs. As illustrated in Figure 2, the inputs encompass a variety of sources and events: e.g., messages, configuration changes, timers. Other less explicit inputs are events such as node failures and random choices.

**How to explore node behavior** Because node behavior depends on the inputs, we want to explore a node’s behavior by subjecting the node to a variety of possible inputs in a way that systematically exercises its code paths. The literature presents us with a technique for this purpose. In software testing, symbolic execution [13] is a technique that explores all possible code paths in a program — symbolic execution treats the input variables of the program as symbolic inputs, and during execution collects the constraints that describe which input values can lead to a particular point in the code. Albeit powerful, this technique comes with significant program execution overhead and, more problematically, it does not easily interact with the environment due to the abstraction of symbolic values.

<sup>1</sup>This is the common approach for model checking distributed systems (e.g., see [31, 43]).

As these two aspects are crucial for testing a real system that runs over the network, we look at a variant of symbolic execution called concolic execution [12, 18, 24] which easily interacts with the environment and has less overhead. This technique executes the code with concrete inputs, while still collecting constraints along code paths. To drive execution down a particular path, the concolic execution engine picks one constraint (e.g., branch predicate) and queries the satisfiability solver to choose a concrete input that negates the constraint.

To exhaustively explore node behavior, we would ideally explore all possible paths at the each node. While concolic execution is in theory capable of exploring all possible code paths, in practice it is severely limited as the number of paths to explore grows exponentially with the size and number of the inputs and the number of branches in the code. We discuss later in this section our insights for dealing with this problem.

**Exercising node behavior to test system states** The node serving as the explorer, runs a *controller*, and a *concolic execution engine*. The controller starts by demanding the creation of a shadow snapshot (Section 3 describes the mechanism we adopt in our prototype). It then uses a previously encountered real input (e.g., a message) to record the constraints encountered on the code path executed with that input (e.g., by invoking a message handler). This initial set of constraints are then passed on to the concolic engine. After completing the initial constraint recording, the concolic engine starts negating predicates one at a time, resulting in a set of inputs, each of which satisfies a particular constraint. To explore a particular input, the controller instructs the explorer’s shadow checkpoint to clone itself, and then resumes running from this cloned checkpoint. The constraints during this new execution path at the explorer node are once again recorded and fed to the concolic engine, which then updates the aggregate set of constraints and keeps producing new inputs. Updating the aggregate set is important for achieving full coverage, since the previous runs might not have reached all branches that exist in the code.

Once the exploration with a particular input completes<sup>2</sup>, the cloned snapshot is checked for faults as explained in the next subsection.

## 2.3 Detecting faults

We detect faults by checking for violations of safety and liveness properties in the cloned snapshots. These properties are user-specified and we assume that they capture system-specific invariants or describe the desired system behavior. Some systems were designed with these types

<sup>2</sup>Deciding on the termination of a distributed computation is a well-understood problem [32].



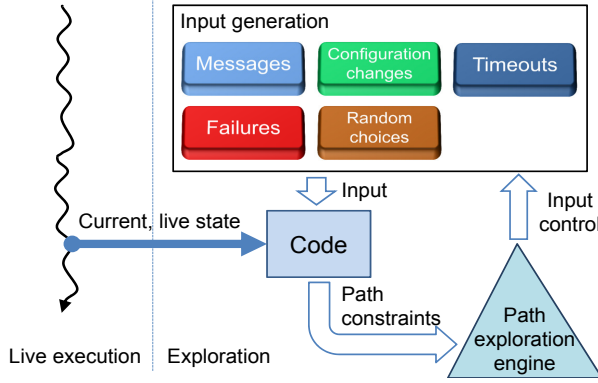


Figure 2: DiCE explores executions from live state, and provides a number of means of controlling inputs that the concolic engine uses to explore system-wide behavior. Exploration occurs in parallel with the production system.

of properties in mind. When that is not the case, the properties can capture the best system practices, as is the case with BGP [26].

**Checking properties across domains** Let  $N$  be the set of nodes, and  $\Theta_i, i \in N$  denote the set of node’s  $i$  states executing in the cloned snapshot. A property, or global check, is expressed as a function  $g(\Theta_1, \Theta_2, \dots, \Theta_{\|N\|}) \in \{0, 1\}$ . Note that a global check considers system-wide behavior and may potentially require accessing information at multiple nodes in different administrative domains.

To control the information shared across domains, we introduce a narrow interface. We consider a subfamily of global checks for which  $g(\Theta_1, \Theta_2, \dots, \Theta_{\|N\|}) = [\sum_{i \in N} f(\Theta_i)] >^? th$ , where  $f(\Theta_i) \in \mathbb{N}_0$  is a check that only accesses local state and  $th$  is a property-specific threshold (e.g., 0).

In this scheme, a centralized entity (e.g., the explorer) computes the global check as the sum of local check values<sup>3</sup>. To preserve privacy, the output of a local check should not contain any private information. For example, local checks can be written in the form: was there a certain change in the node’s state? However, we anticipate there could be cases when individual domains are not willing to disclose local checks unless anonymity can be guaranteed, e.g., if a local check necessarily leaks private information. At the expense of increased computational complexity, we can control information sharing by securely summing local check values so that only the final outcome is known to participating nodes and single addends are not known. Appendix A discusses one such scheme for providing anonymous property checks.

<sup>3</sup>Our scheme requires a global check to be decomposed into local checks, thus this might require an ad-hoc distributed protocol as we show later for detecting policy conflicts.

## 2.4 Choosing symbolic inputs and counter-ing exponential explosion of code paths

We recognize there exists a tension between one’s desire to exhaustively test a system (reflecting a plethora of symbolic inputs) and the hard limits imposed by the exponential explosion of code paths. Here, we present the lessons learned and share our insights in how to find an acceptable balance.

As shown in Figure 2, a variety of means can be used to enable the concolic engine to drive the explorer’s behavior to reach relevant system-wide states. Ideally, we would want to define as symbolic any input that could cause the system to transition from the current state to a faulty state. However, how to identify all these inputs is an open question which might not have a general solution. In our experience, we find that leveraging domain knowledge is an effective approximation. For instance, anticipating some of the discussion from the next section, we identify that a key aspect of DNS name resolution is the *random choice* in querying one of many possible name servers for a given domain name. This driver of node behavior is easy to recognize and, treated as symbolic input, allows to explore interesting interleaving of node actions and to capture the effects of remote *node failures*. Interestingly, property definitions may give hints as to what inputs need to be symbolic. For example, persistent oscillations in BGP can be caused by conflicting policies at different administrative domains. Policies are encoded in router configuration. Treating a policy *configuration change* as symbolic input enables to exercise the BGP route selection process and find potential conflicts lurking in the configuration.

Finally, when dealing with symbolic messages we observe that concolic execution easily ends up creating many invalid inputs that simply exercise the message parsing code. Also, if the message format allows for variable length fields, we note that concolic execution has difficulties in producing messages where such fields are shrunk or grown. Therefore, we find it useful to use grammar-based whitebox-fuzzing [23] which leverages knowledge of the message format to produce a large number of inputs that quickly pass validation checks. We apply the fuzzing code before the message handlers, and rely on the domain knowledge to identify these handlers.

In addition to a thoughtful choice of symbolic inputs, we have two key principles for dealing with the path explosion and large input problems, as listed below:

- Start the exploration from current system state (the shadow snapshot). Doing so eliminates the need to replay from initial state a potentially large history of inputs to reach a desired point in the code.
- Explore behaviors that are a result of small inputs,

both size-wise and in number. The intuition is to try to reach faulty states that are small deviations from current state rather than being more exhaustive with the associated exponential increase in states.

## 2.5 Preventing information leakage

Ideally, the data that is crossing the trust boundary among the nodes should not reveal any confidential information. At a high level, we observe that there are two main kinds of information that can be leaked: potential node behavior and configuration data.

Leakage of node behavior is a direct consequence of systematic code path exploration. We argue however that in a long-running system the behavior has already been revealed for at least the most common set of code paths.

Configuration data can be leaked if the executed code paths produce messages containing a direct copy of the configuration data or an indirect manipulation thereof from which the configuration data can be reverse engineered. However, using concolic execution aids in information hiding. When the concolic engine wants to negate a constraint, it can pick any random value that negates the constraint to drive execution. Thus, the randomized nature of these inputs limits this kind of information leakage. In addition, we can annotate what data is confidential and avoid recording constraints from the code that handles the confidential data so that it cannot leak into the inputs the concolic engine produces.

Finally, additional measures can be taken, including: (i) rate limiting the exploration or responses to property checks, or (ii) refusing certain explorer nodes altogether in the absence of any trust. However, we leave for future work a thorough study of the security-related aspects of our approach.

## 2.6 Discussion

A number of issues, such as the (possibly parallel) order in which the nodes act as explorers, the size of the shadow snapshot, and the amount of resources devoted to exploration at each node are application-specific and orthogonal to this paper; we discuss them in more detail in [1]. Here we only note that it is possible to limit resource consumption during exploration using existing primitives on many platforms.

**Limitations** The types of faults we can detect are a subset of faults that can be detected in a general distributed system [27]. We do not attempt to verify algorithms, protocols, or the operating system of the node. We do not incorporate Byzantine faults. To help it deal with this type of faults, DiCE could directly benefit from schemes that ensure accountability [26]. Further, we only check

for known classes of faults that are captured in the given system-specific properties.

As with any fault detection solution, the potential DiCE issues are false positives and false negatives. DiCE can exhibit false negatives if the given properties are not capable of discerning the faulty state. Specifying liveness properties (which we can check) is helpful in this case as it is easier than exhaustively capturing all relevant safety properties [31]. False negatives also arise when there exists no code path that the concolic engine can exercise with small inputs to reach a faulty state. False positives are less of a problem, as the live execution over the cloned snapshot is evidence of behavior that is the result of processing a particular input. However, the properties themselves should be defined in a way that avoids false positives.

Note that the set of inputs that systematically covers message handling code on one node might not result in full path coverage of other participants (when they run using the inputs they receive in the shadow messages). We cannot easily accommodate system-wide coverage because we would need to share constraints with remote nodes and we deem this unfeasible because of privacy considerations. However, our evaluation with BGP and DNS demonstrates that important classes of faults can be detected without having system-wide path coverage.

Finally, DiCE is not a bug-finding tool that could be used to pinpoint the location of programming errors. DiCE is neither a traditional fault detection tool in the sense that the faults it can find are not detected in the live production system, but rather by reaching faulty states in the cloned snapshots.

## 3 DiCE prototype and applications

This section discusses our DiCE prototype and its application to two federated, heterogeneous distributed systems: BGP and DNS. For each case study, we present a brief overview of the target distributed system and describe how we integrate DiCE with it.

Our prototype consists of a concolic engine, a part written in C and integrated with the target systems, and a Python implementation of the DiCE controller. We use the Oasis [18] concolic engine as the basis for code path exploration. Oasis instruments C programs to record constraints on symbolic inputs during program execution. We discussed in [1] the modifications we made in Oasis. These include support for exploring from current state and the ability to use a single executable where both the original and instrumented code co-exist for avoiding performance overheads in the deployed system while recording constraints during exploration. In addition, in this work we change the Oasis filesystem/network model

to manage shadow connections.

### 3.1 Integration with BGP

Here, we present the case study with BGP and use it to describe the details of our DiCE prototype. We start this section by providing an overview of BGP. We then discuss the integration of DiCE with the BIRD [7] open-source routing daemon. BIRD is written in C and supports multiple routing protocols. It is in production use serving as a route server in several Internet exchange points.

#### 3.1.1 BGP overview

The Internet consists of tens of thousands of domains, so-called autonomous systems (ASes). ASes are typically administered by Internet Service Providers (ISPs). While the ASes have freedom in choosing their intra-domain routing protocol, Border Gateway Protocol (BGP) [37] is the inter-domain routing protocol that acts as the glue that ensures universal connectivity in the Internet and is spoken at each border router.

Each BGP speaker maintains a routing table, or Routing Information Base (RIB) that associates a route to a network prefix with the next hop router and the list of ASes (AS\_PATH) that needs to be taken to reach a given IP in that prefix. The routing information is distilled into a Forwarding Information Base (FIB) that is used to make packet forwarding decisions. BGP speakers establish their routing tables by exchanging UPDATE messages which *announce* routes (each composed of a prefix and a bitmask length) along with their corresponding attributes (e.g., AS\_PATH) and/or *withdraw* routes that are no longer available.

Recently, the protocol has been extended to allow for 4-byte AS numbers [41], and thus the messages can carry the optional AS4\_PATH attribute. Legacy routers that do not understand the 4-byte AS numbers do not attempt to interpret the new attribute and simply pass it along with their updates.

#### 3.1.2 Implementation

For integrating with BIRD, we made a small number of changes that fulfill the application requirements for applying DiCE to BGP. Specifically, (i) we marked the symbolic inputs (only a few LoC), (ii) we added support for taking snapshots and managing shadow connections (about 1300 LoC), and (iii) we exposed certain properties based on the local state that are queried by the controller in order to detect faults (about 200 LoC). We now discuss each of the implementation details.

**Symbolic inputs** For the reasons given in Section 2.4, we choose to treat UPDATE messages and policy config-

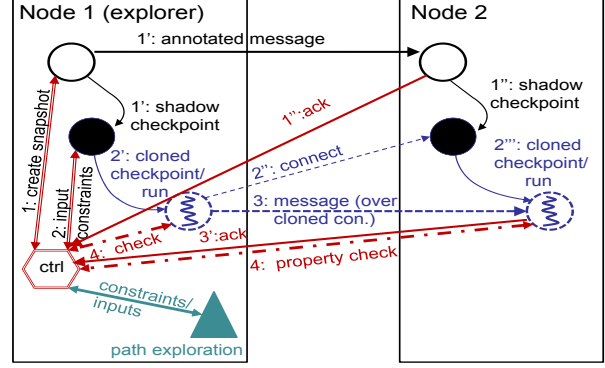


Figure 3: DiCE prototype in action. Thin lines correspond to the checkpoints being created from live state, dashed lines denote cloned state and connections, and the dash-dotted lines are DiCE controller actions.

uration changes as the basis to derive new inputs during exploration.

In BGP, UPDATE messages are the main drivers for state change while the other state changing messages are only responsible for establishing or tearing down peerings and we leave them for future work. As the format of BGP messages is well-defined in the RFC [37], we apply grammar-based fuzzing [23] to the path attributes and we mark the Network Layer Reachability Info (NLRI) region of the message as symbolic. An UPDATE message can carry several path attributes each of which is encoded as a type, length, and value field. To fuzz message attributes, we create two symbolic inputs for each attribute present in the initial message<sup>4</sup>. With respect to the fuzzed message, we assign to these inputs the meaning of attribute presence and length, respectively. In other words, if Oasis picks a non-zero value for the first input we include the attribute, otherwise we remove it from the message; the attribute’s length is matched to the second input. Therefore, Oasis can produce fuzzed messages based on all combinations that these symbolic inputs can have. In addition, Oasis can change the content of the NLRI based on the set of recorded constraints.

We define a further symbolic input that represents changes to a route preference as it would be caused by a configuration change. Specifically, this input reflects for a given route whether that is the most preferred route or not. Thus, Oasis can explore system behaviors for different preferences in the explorer’s route selection process.

**Snapshot** We use Figure 3 as the guiding example of DiCE’s operation in one round of exploration. Recall that the explorer initiates the exploration by triggering a checkpointing phase which results in the creation of a snapshot (Figure 3 step 1). In the current implementation, this is done by taking a checkpoint at the explorer and sending an UPDATE message annotated with a cus-

<sup>4</sup>Except for mandatory attributes as a message without them is an invalid input.

tom path attribute (step 1'). Enclosed in the attribute are the checkpoint number, the IP address of the explorer, a counter that is decreased at each hop to confine the exploration scope, and a weight used for termination detection (explained below).

A router that receives this annotated message interprets the custom attribute and, if this is the first time it sees the current checkpoint number, it takes a checkpoint (step 1''). As BIRD runs a single process, the procedure to take a checkpoint is simply implemented using the `fork` system call. This way of checkpointing allows us to create a large number of checkpoints with a small memory footprint. When a checkpoint is created from the BIRD process running in production, DiCE isolates the forked process from its parent by closing the open sockets<sup>5</sup> and marking them as shadow sockets. Also, the checkpoint is isolated from the FIB. Finally, DiCE opens a new socket to listen for incoming shadow connections on a different port from that used by the production instance of BIRD. With some implementation effort, the same techniques could be applied to other, more complex routing software<sup>6</sup>.

The dissemination of the checkpoint message is achieved by announcing a route to a dedicated prefix so that every router eventually receives the checkpoint-annotated message. However, the explorer needs to be acknowledged when the checkpointing phase ends. For this purpose, we use a variation of the weight-throwing algorithm for termination detection in a distributed system [32]. Briefly, the explorer starts by sending the checkpointing message with an initial weight (e.g., 1). When a router receives the message, it keeps a part of the weight for itself (e.g.,  $weight \cdot 1 / (\#neighbors + 1)$ ) and, while propagating the message, equally shares the remaining part of the weight among its neighbors. A router that does not propagate the message further keeps the received weight for itself. Meanwhile, every router reports its weight to the controller (step 1'''). When the reported weights sum up to the initial weight the controller concludes termination of checkpointing and starts exploring by running the Oasis concolic engine.

**Exploration** Oasis collects constraints along the branches it encounters in the code. In our prototype, the constraints come from: (i) the BIRD C code that deals with UPDATE processing, (ii) the code for fuzzing path attributes, (iii) the code for injecting policy changes, and (iv) the BIRD configuration interpreter. Note that BGP router's behavior is a result of not only the code but also configuration. This is why the concolic engine records constraints for the interpreted configuration. Therefore,

the explored execution paths are comprehensive of both code and configuration.

To perform path exploration, Oasis negates one constraint at a time and produces a new assignment of symbolic inputs (step 2) which are used to drive one *execution* of exploration. First, an isolated BIRD process is forked from the previously established shadow checkpoint (step 2'). Recall we term a process forked from a shadow checkpoint as cloned checkpoint. Before a message exchange between cloned checkpoints can take place, a connection is required to be setup. This is done by connecting<sup>7</sup> to the shadow checkpoint of the message destination (step 2'') which creates a cloned checkpoint (step 2'''). Note that only one cloned checkpoint per node is created for each execution: the first connection is handled by the shadow checkpoint, any subsequent connection is managed by the cloned checkpoint itself. Then, messages are exchanged over these connections (step 3). When it receives the first message, each process ignores the previously existent information about the route(s) contained in the message. This is to ensure that messages are propagated as they would in production, because otherwise the BGP selection would ignore the announcement.

The messages are extended to carry weight information so that the same termination detection algorithm described before is used to detect BGP convergence in the cloned snapshot (step 3'). However, routing may not converge if BGP is in an ill-state [25] within the snapshot. Therefore, during exploration, we use the method in [19] to prevent persistent BGP oscillation under arbitrary filtering (explained in Section 4.3). Lack of convergence due to system dynamics (session failures) are tolerated by shutting down the failed BGP session at the node at which a BGP error occurs.

When the controller detects that one execution terminates, it queries (step 4) all routers that participated in the exploration for properties that allow for fault detection as explained in Section 4. Then, exploration can progress with another execution based on the next negated constraint. When each execution terminates, the processes involved in the exploration can terminate as well and release the resources. The exploration then concludes when Oasis has covered the paths reachable by controlling the composite set of recorded constraints. At the end of the exploration all checkpoint processes are terminated as well.

**Legacy routers and deployment** To capture a system-wide snapshot, the annotated message has to propagate through the network and reach all routers within the exploration scope. This can be easily achieved by reserving a prefix for this purpose which is announced to trigger the

<sup>5</sup>Of course, this does not affect the production instance of BIRD.

<sup>6</sup>For example, Quagga [8] is structured as a set of processes, one per routing protocol. DiCE can be applied by controlling per-protocol shadow connections, and by isolating the processes from the FIB.

<sup>7</sup>This requires a 2-way handshake to avoid race conditions.



checkpointing and withdrawn afterwards. This does not require any modification to BGP because custom route attributes are allowed in the protocol specifications, making it possible to pass-through legacy routers.

We envision that DiCE could be deployed incrementally on BGP routers. To check for faults due to programming errors, an ISP might configure a DiCE-enabled router to send exploration messages to spare equipment which can run in isolation and be monitored for observable errors (e.g., through system logs). In addition, an ISP could check for misconfigurations by deploying a single DiCE-enabled route server configured with the ISP policy and connected with DiCE-enabled routers at the neighboring ISPs.

There is evidence that router manufacturers are starting to leverage the additional computational power; for example, Cisco IOS XR Operating System for core routers is SMP-capable [2].

## 3.2 Integration with DNS

We now build upon the first case study and succinctly describe the important differences for applying DiCE with another crucial system for the Internet infrastructure: Domain Name System (DNS). After a short overview on DNS, we discuss the integration with an open-source DNS server.

### 3.2.1 DNS overview

DNS [33, 34] realizes a name resolution service for the Internet that maps host names to IP addresses. DNS is a distributed database composed of a large number of hierarchically organized, autonomously administered zones, each of which is a subspace of the global namespace that is authoritative for the names that share the same suffix with the zone’s domain name. Each zone maintains a list of so called *Resource Records* (RRs) for the domains under the zone’s authority. For example, the A records map names to IP addresses; the NS records identify authoritative name servers (ANSs).

Typically, name resolution is carried out by a DNS resolver. In the basic form, given a name, the resolver queries one of the ANSs belonging to the name’s domain. If the list of ANSs is not known, the resolver needs first to retrieve it from an ANS of the parent zone. This process can repeat up to the root of the DNS hierarchy that is conveniently hard-coded in all resolvers.

### 3.2.2 Implementation

Using the lessons learned during BGP integration, it took us less than a week to integrate DiCE with MaraDNS [4] 2.0.02, an open-source DNS server. This time includes

all the efforts to compile the codebase, implement lightweight checkpointing, setup an experimental testbed, read the DNS code, decide what to make symbolic, and implement the code that drives exploration. Overall, we added 74 LoC to MaraDNS to integrate with the concolic engine, and another 78 LoC to enable symbolic inputs.

We leverage the fact that DNS servers process queries that do not change their state. This simplifies the integration because the deployed nodes form a snapshot. We only instrument the recursive resolver, and integrate it with DiCE.

**Symbolic inputs** In DNS, local node actions do not result in state changes at remote nodes<sup>8</sup>. In principle, therefore, a single node cannot be responsible for an event like system-wide session resets such as in BGP. However, node behavior is not only driven by code but also by configuration. In the case of DNS, errors lurking in the system configuration are an example of a cause of misbehavior that can be problematic for system reliability (e.g., the impossibility of resolving certain domains [36]). In the absence of state-changing operations, subtle misconfigurations manifest themselves as the result of specific interleaving of node actions. For DNS, that is the particular path (ordering of nodes) in which a DNS resolver attempts to resolve a domain name. Note that this path is also affected by failures of DNS servers or routing instabilities. We therefore recognize the importance of achieving systematic exploration of the system-wide execution paths during DNS resolution.

To drive the exploration, we change the way the resolver decides which ANS to query when it has multiple choices. We introduce a `get_server()` function that for each ANS list, maintains a subset of active servers. Each time the resolver needs a server from that list, the function selects one from the active subset. Both the subset and the server selection are implemented using boolean symbolic inputs. This way, the concolic engine tries all the possible server subsets and all the possible server combinations. In doing so, it mimics the remote server failures that could cause different local choices, as well as the different random choices in choosing a server.

## 4 Evaluation

In this section, we describe the way we evaluated our DiCE prototype, including a detailed description of the properties we use to detect faults.

### 4.1 BGP experimental setup

In our experiments, we make use of a 48-core machine with 64 GB of RAM, running Linux 2.6.30. We install

---

<sup>8</sup>We do not consider security exploits.

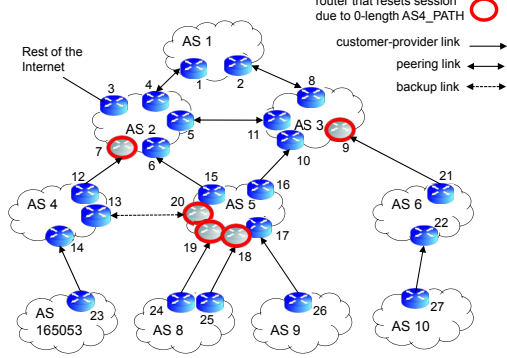


Figure 4: The topology we used in our BGP experiments with harmful global events.

virtual interfaces on this machine, and use them to configure and run multiple BIRD router instances. We have previously quantified the memory overhead (37%) and performance impact (8% in a stress test; negligible in normal operation) of using a concolic engine for path exploration of a BGP router [1].

The 27-router topology we use is shown in Figure 4, and it corresponds to the topology used in [26]. We further annotate the topology with link latencies (30 ms among the routers in the same AS, 5 ms otherwise<sup>9</sup>) and link capacities (620 Mbps), and install it into the ModelNet [39] network emulator running FreeBSD 4.9 on a separate machine. This setup allows us to subject individual packets to link latencies and queuing delays that mimic Internet conditions.

The BIRD instances within an AS are connected in a full-mesh IBGP (internal BGP) topology, while the inter-domain protocol is EBGP (external BGP). This setup does not use route reflectors. The network includes a mix of AS types (tier-1 ASes, tier-2 ASes and small stub ASes) interconnected with either customer-provider, peering or backup relationships as indicated by the arrows in Figure 4. We loaded 319,355 prefixes into the topology by replaying a BGP trace obtained from RouteViews (RIB dump plus 15-min updates trace from routeviews.eqix at April 1, 2010, 17:28 UTC). BGP policies were configured as in [26]: if a route is imported from a customer, it is exported to all neighbors; if a route is from a peer or provider, it is exported to customers only. This set of policies adheres to the Gao-Rexford conditions for stable routing [22] and therefore policy-induced oscillations are not possible. In our network, all ASes except AS 3 use customer route filtering to prevent their customers (including customers’ customers and so on) from injecting advertisements for prefixes they do not own. We already showed in [1] how concolic execution can exercise a BGP router’s behavior in a way that locally

exposes potential origin misconfigurations (route leaks). Because of space limit, we do not present its extension to a system-wide property.

## 4.2 Harmful global events in BGP

An important property we install is the one that checks for the presence of a particular event at a global scale, akin to emergent behavior. An example that affected BGP is the session reset problem [6]. The core of the problem is the fact that the affected routers had difficulty handling an update in which the AS4\_PATH attribute had zero (0) length. The router receiving such an update would not crash, but it would reset the session with the sender of the message. In strict isolation, this seemingly valid handling of a semantically confusing message would not have a far-reaching impact. Unfortunately, a large fraction of routers were not affected by this programming error and were effectively multicasting the session reset signal. Each session reset is followed by a new session establishment that triggers a full routing table download and route processing that is CPU intensive. Moreover, the routing updates containing the confusing AS4\_PATH attributes would then be redelivered to the affected routers, causing another round of session resets, and so on. As a result, the peak update traffic in the Internet was increased by more than a factor of 10 (1000%) [6]. Operators had to manually install packet filters to prevent this fault from recurring.

To enable DiCE to detect this type of a fault, we install the property which signals a possible fault whenever a router resets the session (seen as an increase in the BGP error count) in response to an exploratory message. To trigger this fault in our testbed, we replicate the previously described scenario [6] to the extent possible as we do not have access to the Cisco IOS code. We introduce code into BIRD (already 4-byte AS-compliant) that resets the session when a zero-length AS4\_PATH attributes arrives. The routers that are configured to be affected by the confusing attributes are marked with a circle in Figure 4. Note that AS 165053 is using a 4-byte AS number. The update containing the zero-length AS4\_PATH was generated using a fuzzed message.

**Detection results** We instructed the routers in our testbed to perform exploration after finishing loading the 319,355 prefixes. The routers use an actual message to record the constraints. Different routers explore a different number of iterations on the same code because of the: 1) different messages that end up being used to initially record the constraints, and 2) different C code that gets instantiated by BIRD to enforce the filtering commands. However, Oasis reported that it had explored all paths at each router. The maximum number of explorations was 2002, minimum was 7 while the average number was 763

<sup>9</sup>This mimics the 60 ms RTT across the continental US.

and std. dev. of 586. Routers 23 and 27 explored with 7 inputs because they only accept a default route and have no filtering enabled. The maximum observed time was 670 s to explore the total of 1156 explorations. The average time to explore was 243 s with a std. dev. of 204 s. We also measured the exploration time without accounting for network delays by repeating the experiments with the same initial messages, but without Modelnet. Overall, the times are smaller with an average of 155 s and std. dev. of 113 s.

Given the timescale over which the Internet incidents occurred due the erroneous configuration files and software (likely to be in place for weeks if not months), the time DiCE took to detect these faults is negligible.

**Benefits of using DiCE** Armed with a property that checks the BGP error count (which can be increased due to a variety of different reasons), DiCE produces a list of possible actions that can cause the systemwide error count to go over the threshold. Each ISP would benefit from this advance warning and could take a number of actions, including: 1) notifying the router vendor and requesting a patch, 2) manually fixing the code if the source is available, and 3) installing filters to filter out the offending message (if the action is caused by message). Without DiCE, this kind of repair was undertaken only after the session reset incident took place across the Internet and was diagnosed after several hours [6].

### 4.3 Policy conflicts in BGP

BGP has evolved over time to allow each ISP to independently decide on the set of routes that will be announced to each neighboring AS using a set of *policies*. These policies capture business decisions and are often private. However, conflicting policies can cause undesired persistent routing oscillations [40] which negatively impact end-to-end performance of the Internet.

Note that policy conflicts are due to a design flaw in BGP and only by changing the protocol itself the problem can be definitely addressed. In fact, a recent approach for resolving policy conflicts advocates changing BGP [19]. Here we describe our approach for detecting policy conflicts in a way that does not require protocol modifications for the production traffic. Detecting conflicts is important, for example, to avoid oscillations due to the use of backup routes that were not checked.

Griffin *et al.* [25] formally analyzed the Stable Paths Problem (SPP) which is an abstraction of the problem that BGP intends to solve in a distributed fashion. They attribute policy-inflicted oscillations to a circular set of conflicting rankings between various nodes which form a *dispute wheel*. They show that the absence of a dispute wheel ensures the existence of a unique, stable and robust SPP solution. The work in [22] suggests that

ISPs observe a set of best practices to avoid the dispute wheel, but these are unnecessarily restrictive and difficult to check in a distributed fashion. In order to prevent policy-induced oscillations from occurring, Ee *et al.* [19] augment BGP with the *global precedence value*, which is carried as an additional attribute in route announcements and is used as discriminator in the BGP decision process with higher importance than local policy preferences. The global precedence value is increased by a router whenever the selected route is not the most preferred one according to the local policy. Global precedence values greater than zero exist when routing converges if and only if dispute wheels causing oscillations exist. Put another way, these protocol changes allow a node to break the dispute wheel and stop the oscillations.

The key idea for detecting policy conflicts is to leverage the global precedence value to detect a policy conflict in a cloned snapshot while exploring a large number of possible behaviors. Specifically, we implemented the global precedence value [19] and we used it to ensure that the routing protocol converges<sup>10</sup>. If during convergence a route announcement contains a non-zero global precedence value, it means that the snapshot contains a dispute wheel, and therefore a policy conflict exists.

For our experiments with policy conflicts, we construct a 5-node topology presented in [25], known as GOOD GADGET. This topology presents a Stable Paths solution. However, a single switch in the ranking of paths (policy change) transforms it into a BAD GADGET topology that has a dispute wheel.

DiCE successfully detected the possibility of a policy conflict in this topology by systematically exploring the consequences of one change at a time in the route preference assignments. Overall, there were 75 iterations that took 39 seconds to explore.

**Benefits of using DiCE** The end result of a DiCE run is a list of changes to the explorer’s policy that are bound to cause a conflict. The ability to detect policy conflicts (and the resulting routing oscillations) before they happen on the Internet is beneficial for allowing the freedom of policy decisions in order to accommodate the complex objectives that govern route choices for ISPs.

### 4.4 Cyclic zone dependencies in DNS

Pappas *et al.*, [36] identified three important classes of configuration errors in DNS, and the one they report being particularly difficult to identify is the cyclic zone dependency. This error involves configuration of multiple servers, and importantly, cannot be detected by inspecting the individual server configurations. The error is also insidious in that in the normal course of operation the

<sup>10</sup>Outside of exploration, the usual BGP decision process runs with the node’s policy and unmodified protocol messages.

system functions correctly notwithstanding possible delays. However, a particular failure pattern of authoritative servers can cause resolution to take place via other servers in a domain’s configuration, which then leads to a cyclic dependency involving two or more servers that cannot be resolved. This error results in domain names becoming unresolvable and ultimately unavailable.

We run 5 nodes in the DNS testbed: one recursive DNS server and four authoritative ones. The recursive one uses the “deadwood” resolver from the MaraDNS software package that we integrated with DiCE. Other nodes run standard “maradns” servers, without any source changes. The nodes and their roles in our DNS experiments are as follows:

- Client: wants to determine the IP address of `foo.dd.aaa` and knows only Node1.
- Node1: recursive server. Knows Node2 and Node5.
- Node2: authoritative server for the `aaa` zone. Publishes the following NS records for `dd.aaa`:
  1. `ns.dd.aaa` (Node3)
  2. `ns2.dd.aaa` (Node4)
  3. `ns3.dd.aaa` (unreachable server)
  4. `ns.dd.bbb` (glueless record)
- Node3: authoritative server for the `dd.aaa` zone. Has the IP address of `foo.dd.aaa`
- Node4: redundant authoritative server for the `dd.aaa` zone. Has the IP address of `foo.dd.aaa`
- Node5: authoritative server for the `dd.bbb` zone. Publishes one NS record:
  1. `ns.dd.aaa` (glueless record)

The glueless record is a record that spans different domains without providing “glue” (similar to the A record) in the form of (name, IP address) that could be used to reach a name server.

DiCE successfully detected a cyclic dependency in our testbed after executing 502 explorations that took 532 seconds. The cyclic dependency is locally detected at the resolver. In this case, the `get_server()` function chooses an execution path in which Node3 and Node4 are not queried (e.g., because they are considered to have failed). Query resolution proceeds via Node5, but unfortunately Node5 redirects the query back to Node2, where the same decision to avoid Node3 and Node4 is made again and again.

The cyclic dependency would manifest itself in the production system if Node3 and Node4 were both to fail. This experiment demonstrates how DiCE systematically explores system behavior under possible failures in a case when it is not possible, or difficult, to cause these failures to occur in the production system.

## 5 Related Work

**Model checking.** CrystalBall [42], and MODIST [43] represent the state-of-the-art in model checking distributed system implementations. CrystalBall [42] proactively predicts inconsistencies that can occur in a running distributed system due to unknown programming errors, and effectively prevents them. It works for systems implemented in the Mace [30] framework. CrystalBall nodes periodically collect a consistent snapshot of system state, and locally run a model checking heuristic on the set of state machines instantiated from the snapshot. MODIST [43] is capable of model checking unmodified distributed systems. One could use MODIST to orchestrate state space exploration across a cluster of machines in an isolated (non-deployed) scenario.

DiCE goes beyond these approaches in several important aspects because it: 1) can uncover faults due to inputs that are different than those fed by the model checking harness, 2) deals with the issues arising from federation (need for privacy, inability to retrieve state and configuration), and 3) incorporates the intrinsic heterogeneity of the system (nodes behave differently either due to different implementations, patch-levels, or configurations). Finally, our DiCE prototype can check liveness properties (e.g., BGP convergence), while CrystalBall checks only safety properties.

**Symbolic execution.** Symbolic [13] and concolic execution [12, 18, 24] are effective in discovering bugs in single-machine code by trying to achieve complete coverage of possible code paths. One can use any of these execution engines to look for bugs in distributed systems code. However, they are limited in their ability to reach faulty states as they cannot handle large inputs in long-running systems and realistic configuration files (e.g., Klee [13] can work only with several bytes of input to achieve good path coverage). In addition, these engines are only successful in searching for violations of local assertions (e.g., memory violations). Thus, without the spatial awareness achieved by DiCE, it is not possible to judge the system-wide impact of node actions.

KleeNet [38] builds a test harness that accommodates messaging and fault injection on top of Klee [13]. To search for bugs, KleeNet arranges for path exploration among the set of TinyOs nodes running in isolation on one machine, prior to deployment. This approach is thus similar in spirit to model checking that starts from initial state, along with the shortcomings in dealing with long-running, federated, and heterogeneous systems.

Relative to symbolic execution approaches, DiCE: 1) explores system behavior starting from live state and configuration which is crucial for overcoming the path explosion problem in long-running systems, 2) provides a way to control inputs to a single node that explore rel-



evant federated system states, 3) adapts to the federated environments by providing a narrow interface for sharing the information of local checks, and 4) accommodates system heterogeneity by allowing each administrative domain to separately integrate DiCE.

There have been proposals for performing path exploration at selected points in time on a single-machine [16]. Others have highlighted challenges of fault detection in federated, heterogeneous systems [14]. DiCE is a system that addresses these challenges. In our short paper [1], we present a preliminary DiCE design, and detail our experiences in integrating a BGP router with a concolic execution engine. This paper goes beyond our short paper in that it: 1) shows how to carefully checkpoint the system to allow the concolic engine to extend its reach across the network, 2) shows how to control the inputs to the concolic engine to enable it to reach relevant federated system states, 3) includes a privacy-preserving scheme for checking properties, 4) presents a disjoint set of experimental results involving BGP, and 5) details our experience of integrating DiCE with DNS, along with the accompanying experimental results.

Castro *et al.*, [15] use symbolic execution and constraint solving to randomize inputs that cause application crashes in an effort to improve privacy of bug reports. DiCE applies a similar idea, but to a different domain and for new functionality (fault detection).

#### Application-specific fault detection and prevention

Tools that look for faults in the set of router configurations using static analysis [20] can be quite effective, but cannot check live state spanning multiple nodes, and their configuration (which can differ from the statically checked files). The work on NetReview [26] posits that is difficult to prevent all classes of faults, and argues that the best we can do for the general case is to detect faults in BGP after they occur. DiCE goes one step further in that it detects important classes of faults *before* they manifest themselves.

Alimi *et al.* advocate use of shadow configuration as a network management primitive [9]. This approach installs an alternative configuration within a single ISP’s routers, and checks its validity. DiCE’s shadow snapshot bears resemblance to this “shadow config” primitive, but: 1) is lightweight (works on existing router processes), 2) is automatically created without operator involvement post-deployment, 3) can span multiple ISPs, and 4) serves to detect faults due to unanticipated inputs [6], bugs, or operator mistakes before they are tried out or put into effect. DiCE can benefit from virtual network substrates (e.g., [11]) to simplify shadow and cloned snapshot creation.

Bug-Tolerant Routers [28] run multiple router implementations in parallel using virtualization, and mask faults by voting and changing the environment of the

router processes. In contrast, DiCE possesses the necessary spatial awareness to detect semantic faults that span multiple routers, systematically explores node behavior, and does not require multiple router implementations.

Proposals exist for dealing with specific BGP faults, e.g., oscillations [19]. We argue that it is better to detect a large class of faults before they occur. It is then possible to devise general or specific solutions for preventing them. Our work is complementary to the numerous security extensions to BGP (e.g., [29]) which prevent certain classes of attacks. However, these works cannot guard against programming errors or policy conflicts.

Pappas *et al.*, [35] have proposed and implemented a third-party tool that periodically downloads DNS resource records belonging to a large number of domains, and checks them for cyclic dependencies (as well as other misconfigurations). We demonstrate DiCE’s ability to automatically accomplish a similar task within DNS itself, where there is a clear incentive for the DNS administrators to identify and eliminate cyclic dependencies.

## 6 Conclusions

We presented the design and implementation of DiCE, a system for detecting faults in the long-running, heterogeneous, and federated distributed systems. DiCE enables system operators to first specify properties that capture the desired system behavior. DiCE then: 1) automatically and systematically explores a large number of relevant executions, 2) checks their system-wide impact in isolation while respecting privacy among different administrative entities, and 3) reports property violations. We integrated DiCE with two systems crucial for Internet’s operation: BGP and DNS. This paper describes the lessons we learned on how to control inputs fed to nodes in order to explore relevant system-wide state.

Our evaluation demonstrates DiCE’s effectiveness and ease of integration with existing software written in C. Specifically, our prototype quickly detects faults that can occur due to programming faults, policy conflicts, and misconfigurations. Ultimately, we want to release the toolkit as open source.

## 7 Acknowledgments

We are grateful to Jennifer Rexford, Katerina Argyraki and Jon Crowcroft for their feedback on earlier drafts of this work. This research is funded by a grant from the European Research Council (grant number 259110). Dejan Novaković is supported by the Swiss NSF (grant FNS 200021-125140).

## References

- [1] anonymized for double-blind reviewing.
- [2] Cisco IOS XR Software. <http://www.cisco.com/en/US/products/ps5845/>.
- [3] How To Build A Cybernuke. <http://www.renesys.com/blog/2010/04/how-to-build-a-cybernuke.shtml>.
- [4] MaraDNS. <http://www.maradns.org>.
- [5] Research experiment disrupts Internet, for some. [http://www.computerworld.com/s/article/9182558/Research\\_experiment\\_disrupts\\_Internet\\_for\\_some](http://www.computerworld.com/s/article/9182558/Research_experiment_disrupts_Internet_for_some).
- [6] Staring Into The Gorge: Router Exploits. <http://www.renesys.com/blog/2009/08/staring-into-the-gorge.shtml>.
- [7] The BIRD Internet Routing Daemon. <http://bird.network.cz>.
- [8] The Quagga Routing Suite. <http://www.quagga.net>.
- [9] R. Alimi, Y. Wang, and Y. R. Yang. Shadow Configuration as a Network Management Primitive. In *SIGCOMM*, 2008.
- [10] L. Amini, A. Shaikh, and H. Schulzrinne. Effective Peering for Multi-provider Content Delivery Services. In *INFOCOM*, 2004.
- [11] A. Bavier, N. Feamster, M. Huang, L. Peterson, and J. Rexford. In VINI Veritas: Realistic and Controlled Network Experimentation. In *SIGCOMM*, 2006.
- [12] J. Burnim and K. Sen. Heuristics for Scalable Dynamic Test Generation. Technical Report UCB/EECS-2008-123, EECS Department, University of California, Berkeley, Sep 2008.
- [13] C. Cadar, D. Dunbar, and D. R. Engler. Klee: Unassisted and automatic generation of high-coverage tests for complex systems programs. In *OSDI*, pages 209–224, 2008.
- [14] M. Canini, D. Novaković, V. Jovanović, and D. Kostić. Fault Prediction in Distributed Systems Gone Wild. In *LADIS*, 2010.
- [15] M. Castro, M. Costa, and J.-P. Martin. Better Bug Reporting With Better Privacy. In *ASPLOS*, 2008.
- [16] V. Chipounov, V. Kuznetsov, and G. Candea. S2E: A Platform for In-Vivo Multi-Path Analysis of Software Systems. In *ASPLOS*, 2011.
- [17] B. Cohen. Incentives Build Robustness in BitTorrent. In *P2PECON*, 2003.
- [18] O. Crameri, R. Bachwani, T. Brecht, R. Bianchini, D. Kostić, and W. Zwaenepoel. Oasis: Concolic Execution Driven by Test Suites and Code Modifications. Technical report, EPFL, 2009.
- [19] C. T. Ee, V. Ramachandran, B.-G. Chun, K. Lakshminarayanan, and S. Shenker. Resolving Inter-Domain Policy Disputes. In *SIGCOMM*, 2007.
- [20] N. Feamster and H. Balakrishnan. Detecting BGP Configuration Faults with Static Analysis. In *NSDI*, 2005.
- [21] P.-A. Fouque, G. Poupard, and J. Stern. Sharing Decryption in the Context of Voting or Lotteries. In *FC*, 2001.
- [22] L. Gao and J. Rexford. Stable Internet Routing Without Global Coordination. *IEEE/ACM Trans. Netw.*, 9(6):681–692, 2001.
- [23] P. Godefroid, A. Kiezun, and M. Y. Levin. Grammar-based Whitebox Fuzzing. In *PLDI*, 2008.
- [24] P. Godefroid, N. Klarlund, and K. Sen. DART: Directed Automated Random Testing. In *PLDI*, 2005.
- [25] T. G. Griffin, F. B. Shepherd, and G. Wilfong. The Stable Paths Problem and Interdomain Routing. *IEEE/ACM Trans. Netw.*, 10(2):232–243, 2002.
- [26] A. Haeberlen, I. Avramopoulos, J. Rexford, and P. Druschel. NetReview: Detecting when Interdomain Routing Goes Wrong. In *NSDI*, 2009.
- [27] A. Haeberlen and P. Kuznetsov. The Fault Detection Problem. In *OPODIS*, 2009.
- [28] E. Keller, M. Yu, M. Caesar, and J. Rexford. Virtually Eliminating Router Bugs. In *CoNEXT*, 2009.
- [29] S. Kent, C. Lynn, and K. Seo. Design and Analysis of the Secure Border Gateway Protocol (S-BGP). *DARPA Information Survivability Conference and Exposition*, 2000.
- [30] C. E. Killian, J. W. Anderson, R. Braud, R. Jhala, and A. M. Vahdat. Mace: Language Support for Building Distributed Systems. In *PLDI*, 2007.
- [31] C. E. Killian, J. W. Anderson, R. Jhala, and A. Vahdat. Life, Death, and the Critical Transition: Finding Liveness Bugs in Systems Code. In *NSDI*, 2007.
- [32] A. D. Kshemkalyani and M. Singhal. *Distributed Computing: Principles, Algorithms, and Systems*. Cambridge University Press, 2008.
- [33] P. Mockapetris. Domain Names - Concepts and Facilities IETF RFC 1034. 1987.
- [34] P. Mockapetris. Domain Names - Implementation and Specification IETF RFC 1035. 1987.
- [35] V. Pappas, P. Fältström, D. Massey, and L. Zhang. Distributed DNS Troubleshooting. In *NerT*, 2004.
- [36] V. Pappas, D. Wessels, D. Massey, S. Lu, A. Terzis, and L. Zhang. Impact of Configuration Errors on DNS Robustness. *IEEE J. Sel. A. Commun.*, 27:275–290, 2009.
- [37] Y. Rekhter, T. Li, and S. Hares. A Border Gateway Protocol 4 (BGP-4) IETF RFC 4271. 2006.
- [38] R. Sasnauskas, O. Landsiedel, M. H. Alizai, C. Weise, S. Kowalewski, and K. Wehrle. KleeNet: Discovering Insidious Interaction Bugs in Wireless Sensor Networks Before Deployment. In *IPSN*, 2010.
- [39] A. Vahdat, K. Yocum, K. Walsh, P. Mahadevan, D. Kostić, J. Chase, and D. Becker. Scalability and Accuracy in a Large-Scale Network Emulator. In *OSDI*, 2002.
- [40] K. Varadhan, R. Govindan, and D. Estrin. Persistent Route Oscillations in Inter-Domain Routing. *Computer Networks*, 1996.
- [41] Q. Vohra and E. Chen. BGP Support for Four-octet AS Number Space IETF RFC 4893. 2007.
- [42] M. Yabandeh, N. Knežević, D. Kostić, and V. Kuncak. CrystalBall: Predicting and Preventing Inconsistencies in Deployed Distributed Systems. In *NSDI*, 2009.
- [43] J. Yang, T. Chen, M. Wu, Z. Xu, X. Liu, H. Lin, M. Yang, F. Long, L. Zhang, and L. Zhou. MODIST: Transparent Model Checking of Unmodified Distributed Systems. In *NSDI*, 2009.

## A Anonymous property checks

We now describe a proof-of-concept scheme that uses a protocol for Secure Multy-Party Computation (SMPC) based on the threshold variant of Paillier’s cryptosystem in [21]. Let  $D$  denote the set of participating domains,  $N_j$  be the nodes of domain  $j \in D$ . We assume there exists an out-of-band mechanism for disseminating a shared public key  $PK$  and a list of private keys  $SK_1, \dots, SK_{||D||}$ . Each domain  $j$  sends the cyphertext  $E_{PK}(\sum_{i \in N_j} [f(\Theta_i)])$  to all other domains.

Next, each domain leverages the homomorphic property of the cryptosystem [21] to compute  $c = E_{PK}(\sum_{i \in N} [f(\Theta_i)]) = \prod_{j \in D} E_{PK}(\sum_{i \in N_j} [f(\Theta_i)])$ . The decryption of  $c$  is shared across all domains. Specifically, each domain  $j$  runs a decryption algorithm using  $SK_j$  that produces a decryption share  $c_j$  and sends it to other domains. Finally, each domain inputs  $c_j, \forall j$  to a combiner algorithm that outputs  $\sum_{i \in N} [f(\Theta_i)]$ . Comparing this value with the threshold  $th$  gives the global check.

We implemented the above protocol in Java using “thep”<sup>11</sup> as a starting point. We ran a micro-benchmark to evaluate its performance using the same experimental setup used for BGP. With respect to the experimental topology (Figure 4), we only used one node per AS because we assume that nodes inside the same domain would trust each other. In summary, we obtained that the times needed for one secure computation are 417 ms and 1979 ms for running without and with ModelNet, respectively. This result leads us to a conclusion that a version of DiCE prototype supporting SMPC should implement secure computations as a pipeline running in parallel to the system exploration process, and batch multiple computations together.

---

<sup>11</sup>A Java implementation of Paillier’s cryptosystem <http://code.google.com/p/thep/>.